

regression model using the deep learning Keras library

April 28, 2020

1 Regression model using the deep learning Keras library

Objective: I am going to build a regression model using the Keras library to model the data about concrete compressive strength. -The predictors in the data of concrete strength include:

1. Cement
2. Blast Furnace Slag
3. Fly Ash
4. Water
5. Superplasticizer
6. Coarse Aggregate
7. Fine Aggregate

2 Download and Clean Dataset

Let's start by importing the pandas and the Numpy libraries.

```
[4]: import pandas as pd
import numpy as np
```

```
[5]: concrete_data = pd.read_csv('https://s3-api.us-gio.objectstorage.softlayer.net/
↳cf-courses-data/CognitiveClass/DL0101EN/labs/data/concrete_data.csv')
concrete_data.head()
```

```
[5]:
```

| | Cement | Blast Furnace Slag | Fly Ash | Water | Superplasticizer | \ |
|---|--------|--------------------|---------|-------|------------------|---|
| 0 | 540.0 | 0.0 | 0.0 | 162.0 | 2.5 | |
| 1 | 540.0 | 0.0 | 0.0 | 162.0 | 2.5 | |
| 2 | 332.5 | 142.5 | 0.0 | 228.0 | 0.0 | |
| 3 | 332.5 | 142.5 | 0.0 | 228.0 | 0.0 | |
| 4 | 198.6 | 132.4 | 0.0 | 192.0 | 0.0 | |

| | Coarse Aggregate | Fine Aggregate | Age | Strength |
|---|------------------|----------------|-----|----------|
| 0 | 1040.0 | 676.0 | 28 | 79.99 |
| 1 | 1055.0 | 676.0 | 28 | 61.89 |
| 2 | 932.0 | 594.0 | 270 | 40.27 |
| 3 | 932.0 | 594.0 | 365 | 41.05 |
| 4 | 978.4 | 825.5 | 360 | 44.30 |

```
[6]: concrete_data.shape #data points
```

```
[6]: (1030, 9)
```

```
[7]: concrete_data.describe()
```

```
[7]:
```

| | Cement | Blast Furnace Slag | Fly Ash | Water | \ |
|-------|-------------|--------------------|-------------|-------------|---|
| count | 1030.000000 | 1030.000000 | 1030.000000 | 1030.000000 | |
| mean | 281.167864 | 73.895825 | 54.188350 | 181.567282 | |
| std | 104.506364 | 86.279342 | 63.997004 | 21.354219 | |
| min | 102.000000 | 0.000000 | 0.000000 | 121.800000 | |
| 25% | 192.375000 | 0.000000 | 0.000000 | 164.900000 | |
| 50% | 272.900000 | 22.000000 | 0.000000 | 185.000000 | |
| 75% | 350.000000 | 142.950000 | 118.300000 | 192.000000 | |
| max | 540.000000 | 359.400000 | 200.100000 | 247.000000 | |

| | Superplasticizer | Coarse Aggregate | Fine Aggregate | Age | \ |
|-------|------------------|------------------|----------------|-------------|---|
| count | 1030.000000 | 1030.000000 | 1030.000000 | 1030.000000 | |
| mean | 6.204660 | 972.918932 | 773.580485 | 45.662136 | |
| std | 5.973841 | 77.753954 | 80.175980 | 63.169912 | |
| min | 0.000000 | 801.000000 | 594.000000 | 1.000000 | |
| 25% | 0.000000 | 932.000000 | 730.950000 | 7.000000 | |
| 50% | 6.400000 | 968.000000 | 779.500000 | 28.000000 | |
| 75% | 10.200000 | 1029.400000 | 824.000000 | 56.000000 | |
| max | 32.200000 | 1145.000000 | 992.600000 | 365.000000 | |

| | Strength |
|-------|-------------|
| count | 1030.000000 |
| mean | 35.817961 |
| std | 16.705742 |
| min | 2.330000 |
| 25% | 23.710000 |
| 50% | 34.445000 |
| 75% | 46.135000 |
| max | 82.600000 |

```
[8]: concrete_data.isnull().sum() # CLEAN THE DATA
```

```
[8]: Cement          0
     Blast Furnace Slag  0
     Fly Ash         0
     Water           0
     Superplasticizer  0
     Coarse Aggregate  0
     Fine Aggregate   0
     Age             0
     Strength        0
```

```
dtype: int64
```

2.0.1 Split data into predictors and target

```
[9]: concrete_data_columns = concrete_data.columns

predictors = concrete_data[concrete_data_columns[concrete_data_columns != 'Strength']] # all columns except Strength
target = concrete_data['Strength'] # Strength column
```

```
[10]: predictors.head()
```

```
[10]:
```

| | Cement | Blast Furnace Slag | Fly Ash | Water | Superplasticizer \ |
|---|--------|--------------------|---------|-------|--------------------|
| 0 | 540.0 | 0.0 | 0.0 | 162.0 | 2.5 |
| 1 | 540.0 | 0.0 | 0.0 | 162.0 | 2.5 |
| 2 | 332.5 | 142.5 | 0.0 | 228.0 | 0.0 |
| 3 | 332.5 | 142.5 | 0.0 | 228.0 | 0.0 |
| 4 | 198.6 | 132.4 | 0.0 | 192.0 | 0.0 |

| | Coarse Aggregate | Fine Aggregate | Age |
|---|------------------|----------------|-----|
| 0 | 1040.0 | 676.0 | 28 |
| 1 | 1055.0 | 676.0 | 28 |
| 2 | 932.0 | 594.0 | 270 |
| 3 | 932.0 | 594.0 | 365 |
| 4 | 978.4 | 825.5 | 360 |

```
[11]: target.head()
```

```
[11]:
```

| | |
|---|-------|
| 0 | 79.99 |
| 1 | 61.89 |
| 2 | 40.27 |
| 3 | 41.05 |
| 4 | 44.30 |

Name: Strength, dtype: float64

normalize the data by subtracting the mean and dividing by the standard deviation.

```
[12]: predictors_norm = (predictors - predictors.mean()) / predictors.std()
predictors_norm.head()
```

```
[12]:
```

| | Cement | Blast Furnace Slag | Fly Ash | Water | Superplasticizer \ |
|---|-----------|--------------------|-----------|-----------|--------------------|
| 0 | 2.476712 | -0.856472 | -0.846733 | -0.916319 | -0.620147 |
| 1 | 2.476712 | -0.856472 | -0.846733 | -0.916319 | -0.620147 |
| 2 | 0.491187 | 0.795140 | -0.846733 | 2.174405 | -1.038638 |
| 3 | 0.491187 | 0.795140 | -0.846733 | 2.174405 | -1.038638 |
| 4 | -0.790075 | 0.678079 | -0.846733 | 0.488555 | -1.038638 |

| | Coarse Aggregate | Fine Aggregate | Age |
|---|------------------|----------------|-----------|
| 0 | 0.862735 | -1.217079 | -0.279597 |
| 1 | 1.055651 | -1.217079 | -0.279597 |
| 2 | -0.526262 | -2.239829 | 3.551340 |
| 3 | -0.526262 | -2.239829 | 5.055221 |
| 4 | 0.070492 | 0.647569 | 4.976069 |

2.0.2 save the number of predictors to n_cols

```
[13]: n_cols = predictors_norm.shape[1] # number of predictors
```

3 Import the Keras library

```
[14]: import keras
```

Using TensorFlow backend.

```
/home/jupyterlab/conda/envs/python/lib/python3.6/site-
packages/tensorflow/python/framework/dtypes.py:519: FutureWarning: Passing
(type, 1) or '1type' as a synonym of type is deprecated; in a future version of
numpy, it will be understood as (type, (1,)) / '(1,)type'.
```

```
_np_qint8 = np.dtype(["qint8", np.int8, 1])
```

```
/home/jupyterlab/conda/envs/python/lib/python3.6/site-
packages/tensorflow/python/framework/dtypes.py:520: FutureWarning: Passing
(type, 1) or '1type' as a synonym of type is deprecated; in a future version of
numpy, it will be understood as (type, (1,)) / '(1,)type'.
```

```
_np_quint8 = np.dtype(["quint8", np.uint8, 1])
```

```
/home/jupyterlab/conda/envs/python/lib/python3.6/site-
packages/tensorflow/python/framework/dtypes.py:521: FutureWarning: Passing
(type, 1) or '1type' as a synonym of type is deprecated; in a future version of
numpy, it will be understood as (type, (1,)) / '(1,)type'.
```

```
_np_qint16 = np.dtype(["qint16", np.int16, 1])
```

```
/home/jupyterlab/conda/envs/python/lib/python3.6/site-
packages/tensorflow/python/framework/dtypes.py:522: FutureWarning: Passing
(type, 1) or '1type' as a synonym of type is deprecated; in a future version of
numpy, it will be understood as (type, (1,)) / '(1,)type'.
```

```
_np_quint16 = np.dtype(["quint16", np.uint16, 1])
```

```
/home/jupyterlab/conda/envs/python/lib/python3.6/site-
packages/tensorflow/python/framework/dtypes.py:523: FutureWarning: Passing
(type, 1) or '1type' as a synonym of type is deprecated; in a future version of
numpy, it will be understood as (type, (1,)) / '(1,)type'.
```

```
_np_qint32 = np.dtype(["qint32", np.int32, 1])
```

```
/home/jupyterlab/conda/envs/python/lib/python3.6/site-
packages/tensorflow/python/framework/dtypes.py:528: FutureWarning: Passing
(type, 1) or '1type' as a synonym of type is deprecated; in a future version of
numpy, it will be understood as (type, (1,)) / '(1,)type'.
```

```
np_resource = np.dtype(["resource", np.ubyte, 1])
```

```
[15]: from keras.models import Sequential
      from keras.layers import Dense
```

4 Build a Neural Network

Model that has one hidden layer with 10 neurons and a ReLU activation function. It uses the adam optimizer and the mean squared error as the loss function.

```
[16]: # define regression model
def regression_model():
    # create model
    model = Sequential()
    model.add(Dense(10, activation='relu', input_shape=(n_cols,)))
    model.add(Dense(10, activation='relu'))
    model.add(Dense(1))

    # compile model
    model.compile(optimizer='adam', loss='mean_squared_error')
    return model
```

4.0.1 import scikit-learn in order to randomly split the data into a training and test sets

```
[17]: from sklearn.model_selection import train_test_split
```

4.0.2 Splitting the data into a training and test sets by holding 30% of the data for testing

```
[18]: X_train, X_test, y_train, y_test = train_test_split(predictors, target,
    ↪ test_size=0.3, random_state=42)
```

5 Train and Test the Network

```
[19]: # build the model
      model = regression_model()
```

5.0.1 Next, we will train and test the model at the same time using the fit method. We will leave out 30% of the data for validation and we will train the model for 50 epochs.

```
[20]: # fit the model
      model.fit(predictors_norm, target, validation_split=0.3, epochs=50, verbose=1)
```

Train on 721 samples, validate on 309 samples

Epoch 1/50

721/721 [=====] - 1s 1ms/step - loss: 1726.3008 -

```

val_loss: 1236.9262
Epoch 2/50
721/721 [=====] - 0s 387us/step - loss: 1693.4964 -
val_loss: 1214.2847
Epoch 3/50
721/721 [=====] - 0s 439us/step - loss: 1662.3141 -
val_loss: 1191.7048
Epoch 4/50
721/721 [=====] - 0s 436us/step - loss: 1629.5006 -
val_loss: 1167.0351
Epoch 5/50
721/721 [=====] - 0s 389us/step - loss: 1589.3327 -
val_loss: 1138.3053
Epoch 6/50
721/721 [=====] - 0s 411us/step - loss: 1539.1931 -
val_loss: 1102.9854
Epoch 7/50
721/721 [=====] - 0s 365us/step - loss: 1476.7772 -
val_loss: 1059.5614
Epoch 8/50
721/721 [=====] - 0s 411us/step - loss: 1397.8101 -
val_loss: 1005.5447
Epoch 9/50
721/721 [=====] - 0s 442us/step - loss: 1301.9842 -
val_loss: 942.7596
Epoch 10/50
721/721 [=====] - 0s 384us/step - loss: 1190.5798 -
val_loss: 873.6619
Epoch 11/50
721/721 [=====] - 0s 393us/step - loss: 1070.5728 -
val_loss: 796.3750
Epoch 12/50
721/721 [=====] - 0s 341us/step - loss: 937.8116 -
val_loss: 718.0617
Epoch 13/50
721/721 [=====] - 0s 382us/step - loss: 806.5951 -
val_loss: 637.1517
Epoch 14/50
721/721 [=====] - 0s 442us/step - loss: 680.2037 -
val_loss: 559.3722
Epoch 15/50
721/721 [=====] - 0s 395us/step - loss: 567.0572 -
val_loss: 487.0216
Epoch 16/50
721/721 [=====] - 0s 433us/step - loss: 472.4213 -
val_loss: 424.3209
Epoch 17/50
721/721 [=====] - 0s 446us/step - loss: 399.7081 -

```

```

val_loss: 373.0173
Epoch 18/50
721/721 [=====] - 0s 391us/step - loss: 347.5732 -
val_loss: 329.4477
Epoch 19/50
721/721 [=====] - 0s 397us/step - loss: 311.3485 -
val_loss: 295.9437
Epoch 20/50
721/721 [=====] - 0s 385us/step - loss: 287.4794 -
val_loss: 270.8085
Epoch 21/50
721/721 [=====] - 0s 380us/step - loss: 271.5739 -
val_loss: 251.2509
Epoch 22/50
721/721 [=====] - 0s 387us/step - loss: 260.9291 -
val_loss: 236.1270
Epoch 23/50
721/721 [=====] - 0s 365us/step - loss: 251.8219 -
val_loss: 225.3402
Epoch 24/50
721/721 [=====] - 0s 421us/step - loss: 244.8533 -
val_loss: 216.1450
Epoch 25/50
721/721 [=====] - 0s 446us/step - loss: 238.7277 -
val_loss: 208.2911
Epoch 26/50
721/721 [=====] - 0s 444us/step - loss: 233.4298 -
val_loss: 201.6223
Epoch 27/50
721/721 [=====] - 0s 367us/step - loss: 228.8940 -
val_loss: 195.9712
Epoch 28/50
721/721 [=====] - 0s 415us/step - loss: 224.8551 -
val_loss: 191.7873
Epoch 29/50
721/721 [=====] - 0s 388us/step - loss: 220.9874 -
val_loss: 187.6183
Epoch 30/50
721/721 [=====] - 0s 412us/step - loss: 217.3150 -
val_loss: 184.3843
Epoch 31/50
721/721 [=====] - 0s 466us/step - loss: 213.8209 -
val_loss: 181.7941
Epoch 32/50
721/721 [=====] - 0s 412us/step - loss: 210.7043 -
val_loss: 178.8490
Epoch 33/50
721/721 [=====] - 0s 441us/step - loss: 207.5997 -

```

```

val_loss: 174.9418
Epoch 34/50
721/721 [=====] - 0s 368us/step - loss: 204.7411 -
val_loss: 172.2746
Epoch 35/50
721/721 [=====] - 0s 445us/step - loss: 202.3195 -
val_loss: 169.3990
Epoch 36/50
721/721 [=====] - 0s 417us/step - loss: 199.6437 -
val_loss: 167.7714
Epoch 37/50
721/721 [=====] - 0s 437us/step - loss: 197.7970 -
val_loss: 166.5286
Epoch 38/50
721/721 [=====] - 0s 395us/step - loss: 195.4487 -
val_loss: 163.8821
Epoch 39/50
721/721 [=====] - 0s 412us/step - loss: 193.2727 -
val_loss: 160.8351
Epoch 40/50
721/721 [=====] - 0s 363us/step - loss: 191.1801 -
val_loss: 160.4910
Epoch 41/50
721/721 [=====] - 0s 470us/step - loss: 188.9282 -
val_loss: 157.8214
Epoch 42/50
721/721 [=====] - 0s 415us/step - loss: 186.9619 -
val_loss: 156.2641
Epoch 43/50
721/721 [=====] - 0s 393us/step - loss: 185.1438 -
val_loss: 155.3652
Epoch 44/50
721/721 [=====] - 0s 439us/step - loss: 183.2322 -
val_loss: 154.3583
Epoch 45/50
721/721 [=====] - 0s 446us/step - loss: 181.7427 -
val_loss: 153.6083
Epoch 46/50
721/721 [=====] - 0s 393us/step - loss: 180.0934 -
val_loss: 152.6627
Epoch 47/50
721/721 [=====] - 0s 411us/step - loss: 178.4411 -
val_loss: 151.0056
Epoch 48/50
721/721 [=====] - 0s 475us/step - loss: 177.0964 -
val_loss: 150.1668
Epoch 49/50
721/721 [=====] - 0s 390us/step - loss: 175.5285 -

```



```

val_loss: 149.6706
Epoch 50/50
721/721 [=====] - 0s 493us/step - loss: 173.9959 -
val_loss: 149.3502

```

[20]: <keras.callbacks.History at 0x7f13ad551e80>

6 Evaluate the model on the test data

```

[21]: loss_val = model.evaluate(X_test, y_test)
      y_pred = model.predict(X_test)
      loss_val

```

```

309/309 [=====] - 0s 108us/step

```

[21]: 75354250.45954692

6.0.1 compute the mean squared error between the predicted concrete strength and the actual concrete strength.

Let's import the mean_squared_error function from Scikit-learn.

```

[22]: from sklearn.metrics import mean_squared_error

```

```

[23]: mean_square_error = mean_squared_error(y_test, y_pred)
      mean = np.mean(mean_square_error)
      standard_deviation = np.std(mean_square_error)
      print(mean, standard_deviation)

```

```

75354248.91785064 0.0

```

7 Build a Neural Network

```

[24]: total_mean_squared_errors = 50
      epochs = 100
      mean_squared_errors = []
      for i in range(0, total_mean_squared_errors):
          X_train, X_test, y_train, y_test = train_test_split(predictors_norm,
          ↪target, test_size=0.3, random_state=i)
          model.fit(X_train, y_train, epochs=epochs, verbose=0)
          MSE = model.evaluate(X_test, y_test, verbose=0)
          print("MSE "+str(i+1)+": "+str(MSE))
          y_pred = model.predict(X_test)
          mean_square_error = mean_squared_error(y_test, y_pred)
          mean_squared_errors.append(mean_square_error)

      mean_squared_errors = np.array(mean_squared_errors)

```

```

mean = np.mean(mean_squared_errors)
standard_deviation = np.std(mean_squared_errors)

print('\n')
print("Below is the mean and standard deviation of "␣
↪+str(total_mean_squared_errors) + " mean squared errors with normalized data.
↪ Total number of epochs for each training is: " +str(epochs) + "\n")
print("Mean: "+str(mean))
print("Standard Deviation: "+str(standard_deviation))

```

```

MSE 1: 94.0861078898112
MSE 2: 69.73465504075331
MSE 3: 41.54775376921719
MSE 4: 42.71618091867194
MSE 5: 37.36165171996675
MSE 6: 35.69694925203292
MSE 7: 38.464529660913165
MSE 8: 29.792352201097607
MSE 9: 31.08339798566207
MSE 10: 32.261866942964325
MSE 11: 28.205467051286913
MSE 12: 24.49172276740707
MSE 13: 32.455063603842525
MSE 14: 32.137673263796714
MSE 15: 30.564119271090117
MSE 16: 23.42215183097568
MSE 17: 28.91355815751653
MSE 18: 27.627063269754057
MSE 19: 23.75466105390135
MSE 20: 26.9092894742404
MSE 21: 25.579971757907312
MSE 22: 27.248190574275636
MSE 23: 19.101532939182515
MSE 24: 24.28075494735372
MSE 25: 26.657361755864905
MSE 26: 27.359636991926767
MSE 27: 22.512651622488274
MSE 28: 23.38672015119139
MSE 29: 27.738358229109384
MSE 30: 23.105451330783684
MSE 31: 22.600225936247693
MSE 32: 22.59083691692661
MSE 33: 20.17199847150389
MSE 34: 22.966079261310664
MSE 35: 22.310424730615708
MSE 36: 25.68064350757784
MSE 37: 19.228271114016042

```

MSE 38: 22.293625840862976
MSE 39: 25.54287717257503
MSE 40: 19.012725265280714
MSE 41: 22.582373813518043
MSE 42: 20.411192156350342
MSE 43: 23.11588991961433
MSE 44: 25.820463137333448
MSE 45: 26.075211719401832
MSE 46: 21.465590862780328
MSE 47: 21.12123419011681
MSE 48: 21.25171399656623
MSE 49: 21.18800794731066
MSE 50: 24.3175544923949

Below is the mean and standard deviation of 50 mean squared errors with normalized data. Total number of epochs for each training is: 100

Mean: 28.558875355673162

Standard Deviation: 12.449009876691735

[]: