

regression model using the deep learning Keras library

April 28, 2020

1 Regression model using the deep learning Keras library

Objective: I am going to build a regression model using the Keras library to model the data about concrete compressive strength. -The predictors in the data of concrete strength include:

1. Cement
2. Blast Furnace Slag
3. Fly Ash
4. Water
5. Superplasticizer
6. Coarse Aggregate
7. Fine Aggregate

2 Download and Clean Dataset

Let's start by importing the pandas and the Numpy libraries.

```
[1]: import pandas as pd
import numpy as np
```

```
[3]: concrete_data = pd.read_csv('https://s3-api.us-gio.objectstorage.softlayer.net/
↳cf-courses-data/CognitiveClass/DL0101EN/labs/data/concrete_data.csv')
concrete_data.head()
```

```
[3]:
```

| | Cement | Blast Furnace Slag | Fly Ash | Water | Superplasticizer | \ |
|---|--------|--------------------|---------|-------|------------------|---|
| 0 | 540.0 | 0.0 | 0.0 | 162.0 | 2.5 | |
| 1 | 540.0 | 0.0 | 0.0 | 162.0 | 2.5 | |
| 2 | 332.5 | 142.5 | 0.0 | 228.0 | 0.0 | |
| 3 | 332.5 | 142.5 | 0.0 | 228.0 | 0.0 | |
| 4 | 198.6 | 132.4 | 0.0 | 192.0 | 0.0 | |

| | Coarse Aggregate | Fine Aggregate | Age | Strength |
|---|------------------|----------------|-----|----------|
| 0 | 1040.0 | 676.0 | 28 | 79.99 |
| 1 | 1055.0 | 676.0 | 28 | 61.89 |
| 2 | 932.0 | 594.0 | 270 | 40.27 |
| 3 | 932.0 | 594.0 | 365 | 41.05 |
| 4 | 978.4 | 825.5 | 360 | 44.30 |

```
[6]: concrete_data.shape #data points
```

```
[6]: (1030, 9)
```

```
[5]: concrete_data.describe()
```

```
[5]:
```

| | Cement | Blast Furnace Slag | Fly Ash | Water | \ |
|-------|-------------|--------------------|-------------|-------------|---|
| count | 1030.000000 | 1030.000000 | 1030.000000 | 1030.000000 | |
| mean | 281.167864 | 73.895825 | 54.188350 | 181.567282 | |
| std | 104.506364 | 86.279342 | 63.997004 | 21.354219 | |
| min | 102.000000 | 0.000000 | 0.000000 | 121.800000 | |
| 25% | 192.375000 | 0.000000 | 0.000000 | 164.900000 | |
| 50% | 272.900000 | 22.000000 | 0.000000 | 185.000000 | |
| 75% | 350.000000 | 142.950000 | 118.300000 | 192.000000 | |
| max | 540.000000 | 359.400000 | 200.100000 | 247.000000 | |

| | Superplasticizer | Coarse Aggregate | Fine Aggregate | Age | \ |
|-------|------------------|------------------|----------------|-------------|---|
| count | 1030.000000 | 1030.000000 | 1030.000000 | 1030.000000 | |
| mean | 6.204660 | 972.918932 | 773.580485 | 45.662136 | |
| std | 5.973841 | 77.753954 | 80.175980 | 63.169912 | |
| min | 0.000000 | 801.000000 | 594.000000 | 1.000000 | |
| 25% | 0.000000 | 932.000000 | 730.950000 | 7.000000 | |
| 50% | 6.400000 | 968.000000 | 779.500000 | 28.000000 | |
| 75% | 10.200000 | 1029.400000 | 824.000000 | 56.000000 | |
| max | 32.200000 | 1145.000000 | 992.600000 | 365.000000 | |

| | Strength |
|-------|-------------|
| count | 1030.000000 |
| mean | 35.817961 |
| std | 16.705742 |
| min | 2.330000 |
| 25% | 23.710000 |
| 50% | 34.445000 |
| 75% | 46.135000 |
| max | 82.600000 |

```
[7]: concrete_data.isnull().sum() # CLEAN THE DATA
```

```
[7]: Cement          0
     Blast Furnace Slag  0
     Fly Ash         0
     Water           0
     Superplasticizer  0
     Coarse Aggregate  0
     Fine Aggregate   0
     Age             0
     Strength        0
```

```
dtype: int64
```

2.0.1 Split data into predictors and target

```
[8]: concrete_data_columns = concrete_data.columns

predictors = concrete_data[concrete_data_columns[concrete_data_columns != 'Strength']] # all columns except Strength
target = concrete_data['Strength'] # Strength column
```

```
[9]: predictors.head()
```

```
[9]:
```

| | Cement | Blast Furnace Slag | Fly Ash | Water | Superplasticizer \ |
|---|--------|--------------------|---------|-------|--------------------|
| 0 | 540.0 | 0.0 | 0.0 | 162.0 | 2.5 |
| 1 | 540.0 | 0.0 | 0.0 | 162.0 | 2.5 |
| 2 | 332.5 | 142.5 | 0.0 | 228.0 | 0.0 |
| 3 | 332.5 | 142.5 | 0.0 | 228.0 | 0.0 |
| 4 | 198.6 | 132.4 | 0.0 | 192.0 | 0.0 |

| | Coarse Aggregate | Fine Aggregate | Age |
|---|------------------|----------------|-----|
| 0 | 1040.0 | 676.0 | 28 |
| 1 | 1055.0 | 676.0 | 28 |
| 2 | 932.0 | 594.0 | 270 |
| 3 | 932.0 | 594.0 | 365 |
| 4 | 978.4 | 825.5 | 360 |

```
[10]: target.head()
```

```
[10]:
```

| | |
|---|-------|
| 0 | 79.99 |
| 1 | 61.89 |
| 2 | 40.27 |
| 3 | 41.05 |
| 4 | 44.30 |

Name: Strength, dtype: float64

normalize the data by subtracting the mean and dividing by the standard deviation.

```
[11]: predictors_norm = (predictors - predictors.mean()) / predictors.std()
predictors_norm.head()
```

```
[11]:
```

| | Cement | Blast Furnace Slag | Fly Ash | Water | Superplasticizer \ |
|---|-----------|--------------------|-----------|-----------|--------------------|
| 0 | 2.476712 | -0.856472 | -0.846733 | -0.916319 | -0.620147 |
| 1 | 2.476712 | -0.856472 | -0.846733 | -0.916319 | -0.620147 |
| 2 | 0.491187 | 0.795140 | -0.846733 | 2.174405 | -1.038638 |
| 3 | 0.491187 | 0.795140 | -0.846733 | 2.174405 | -1.038638 |
| 4 | -0.790075 | 0.678079 | -0.846733 | 0.488555 | -1.038638 |

| | Coarse Aggregate | Fine Aggregate | Age |
|---|------------------|----------------|-----------|
| 0 | 0.862735 | -1.217079 | -0.279597 |
| 1 | 1.055651 | -1.217079 | -0.279597 |
| 2 | -0.526262 | -2.239829 | 3.551340 |
| 3 | -0.526262 | -2.239829 | 5.055221 |
| 4 | 0.070492 | 0.647569 | 4.976069 |

2.0.2 save the number of predictors to n_cols

```
[13]: n_cols = predictors_norm.shape[1] # number of predictors
```

3 Import the Keras library

```
[14]: import keras
```

Using TensorFlow backend.

```
/home/jupyterlab/conda/envs/python/lib/python3.6/site-  
packages/tensorflow/python/framework/dtypes.py:519: FutureWarning: Passing  
(type, 1) or '1type' as a synonym of type is deprecated; in a future version of  
numpy, it will be understood as (type, (1,)) / '(1,)type'.
```

```
_np_qint8 = np.dtype(["qint8", np.int8, 1])  
/home/jupyterlab/conda/envs/python/lib/python3.6/site-  
packages/tensorflow/python/framework/dtypes.py:520: FutureWarning: Passing  
(type, 1) or '1type' as a synonym of type is deprecated; in a future version of  
numpy, it will be understood as (type, (1,)) / '(1,)type'.
```

```
_np_quint8 = np.dtype(["quint8", np.uint8, 1])  
/home/jupyterlab/conda/envs/python/lib/python3.6/site-  
packages/tensorflow/python/framework/dtypes.py:521: FutureWarning: Passing  
(type, 1) or '1type' as a synonym of type is deprecated; in a future version of  
numpy, it will be understood as (type, (1,)) / '(1,)type'.
```

```
_np_qint16 = np.dtype(["qint16", np.int16, 1])  
/home/jupyterlab/conda/envs/python/lib/python3.6/site-  
packages/tensorflow/python/framework/dtypes.py:522: FutureWarning: Passing  
(type, 1) or '1type' as a synonym of type is deprecated; in a future version of  
numpy, it will be understood as (type, (1,)) / '(1,)type'.
```

```
_np_quint16 = np.dtype(["quint16", np.uint16, 1])  
/home/jupyterlab/conda/envs/python/lib/python3.6/site-  
packages/tensorflow/python/framework/dtypes.py:523: FutureWarning: Passing  
(type, 1) or '1type' as a synonym of type is deprecated; in a future version of  
numpy, it will be understood as (type, (1,)) / '(1,)type'.
```

```
_np_qint32 = np.dtype(["qint32", np.int32, 1])  
/home/jupyterlab/conda/envs/python/lib/python3.6/site-  
packages/tensorflow/python/framework/dtypes.py:528: FutureWarning: Passing  
(type, 1) or '1type' as a synonym of type is deprecated; in a future version of  
numpy, it will be understood as (type, (1,)) / '(1,)type'.
```

```
np_resource = np.dtype(["resource", np.ubyte, 1])
```

```
[15]: from keras.models import Sequential
      from keras.layers import Dense
```

4 Build a Neural Network

Model that has one hidden layer with 10 neurons and a ReLU activation function. It uses the adam optimizer and the mean squared error as the loss function.

```
[17]: # define regression model
      def regression_model():
          # create model
          model = Sequential()
          model.add(Dense(10, activation='relu', input_shape=(n_cols,)))
          model.add(Dense(10, activation='relu'))
          model.add(Dense(1))

          # compile model
          model.compile(optimizer='adam', loss='mean_squared_error')
          return model
```

4.0.1 import scikit-learn in order to randomly split the data into a training and test sets

```
[18]: from sklearn.model_selection import train_test_split
```

4.0.2 Splitting the data into a training and test sets by holding 30% of the data for testing

```
[20]: X_train, X_test, y_train, y_test = train_test_split(predictors, target,
      ↪test_size=0.3, random_state=42)
```

5 Train and Test the Network

```
[21]: # build the model
      model = regression_model()
```

5.0.1 Next, we will train and test the model at the same time using the fit method. We will leave out 30% of the data for validation and we will train the model for 50 epochs.

```
[23]: # fit the model
      model.fit(predictors_norm, target, validation_split=0.3, epochs=50, verbose=1)
```

Train on 721 samples, validate on 309 samples

Epoch 1/50

721/721 [=====] - 0s 412us/step - loss: 158.8141 -

```

val_loss: 106.1055
Epoch 2/50
721/721 [=====] - 0s 445us/step - loss: 157.2554 -
val_loss: 106.7075
Epoch 3/50
721/721 [=====] - 0s 416us/step - loss: 155.6376 -
val_loss: 107.5418
Epoch 4/50
721/721 [=====] - 0s 357us/step - loss: 154.1690 -
val_loss: 107.9913
Epoch 5/50
721/721 [=====] - 0s 368us/step - loss: 152.7373 -
val_loss: 108.5248
Epoch 6/50
721/721 [=====] - 0s 397us/step - loss: 151.0639 -
val_loss: 109.8685
Epoch 7/50
721/721 [=====] - 0s 437us/step - loss: 149.7963 -
val_loss: 110.5298
Epoch 8/50
721/721 [=====] - 0s 392us/step - loss: 148.3151 -
val_loss: 110.5286
Epoch 9/50
721/721 [=====] - 0s 409us/step - loss: 146.9782 -
val_loss: 111.2455
Epoch 10/50
721/721 [=====] - 0s 361us/step - loss: 145.7472 -
val_loss: 112.8683
Epoch 11/50
721/721 [=====] - 0s 415us/step - loss: 144.3118 -
val_loss: 113.2248
Epoch 12/50
721/721 [=====] - 0s 383us/step - loss: 143.0453 -
val_loss: 114.0196
Epoch 13/50
721/721 [=====] - 0s 442us/step - loss: 142.1584 -
val_loss: 115.1693
Epoch 14/50
721/721 [=====] - 0s 381us/step - loss: 140.8834 -
val_loss: 116.3022
Epoch 15/50
721/721 [=====] - 0s 390us/step - loss: 139.5412 -
val_loss: 116.5498
Epoch 16/50
721/721 [=====] - 0s 383us/step - loss: 138.5614 -
val_loss: 117.9367
Epoch 17/50
721/721 [=====] - 0s 388us/step - loss: 137.1343 -

```

```

val_loss: 118.5434
Epoch 18/50
721/721 [=====] - 0s 385us/step - loss: 136.1445 -
val_loss: 120.3195
Epoch 19/50
721/721 [=====] - 0s 360us/step - loss: 134.9155 -
val_loss: 121.5792
Epoch 20/50
721/721 [=====] - 0s 365us/step - loss: 133.8227 -
val_loss: 121.2784
Epoch 21/50
721/721 [=====] - 0s 359us/step - loss: 132.7879 -
val_loss: 122.6053
Epoch 22/50
721/721 [=====] - 0s 472us/step - loss: 131.7246 -
val_loss: 123.6060
Epoch 23/50
721/721 [=====] - 0s 410us/step - loss: 130.6438 -
val_loss: 124.5070
Epoch 24/50
721/721 [=====] - 0s 522us/step - loss: 129.6374 -
val_loss: 125.9104
Epoch 25/50
721/721 [=====] - 0s 447us/step - loss: 128.9126 -
val_loss: 127.3362
Epoch 26/50
721/721 [=====] - 0s 389us/step - loss: 127.8902 -
val_loss: 128.3615
Epoch 27/50
721/721 [=====] - 0s 438us/step - loss: 127.0193 -
val_loss: 128.7330
Epoch 28/50
721/721 [=====] - 0s 415us/step - loss: 125.9038 -
val_loss: 129.0002
Epoch 29/50
721/721 [=====] - 0s 410us/step - loss: 124.8321 -
val_loss: 129.0912
Epoch 30/50
721/721 [=====] - 0s 417us/step - loss: 123.7985 -
val_loss: 131.1239
Epoch 31/50
721/721 [=====] - 0s 390us/step - loss: 122.9441 -
val_loss: 132.4291
Epoch 32/50
721/721 [=====] - 0s 442us/step - loss: 122.2538 -
val_loss: 133.5568
Epoch 33/50
721/721 [=====] - 0s 390us/step - loss: 121.3691 -

```

```

val_loss: 132.2868
Epoch 34/50
721/721 [=====] - 0s 438us/step - loss: 120.4920 -
val_loss: 134.6746
Epoch 35/50
721/721 [=====] - 0s 473us/step - loss: 119.8096 -
val_loss: 135.7861
Epoch 36/50
721/721 [=====] - 0s 414us/step - loss: 118.9906 -
val_loss: 136.8689
Epoch 37/50
721/721 [=====] - 0s 448us/step - loss: 118.0742 -
val_loss: 136.6626
Epoch 38/50
721/721 [=====] - 0s 382us/step - loss: 117.2895 -
val_loss: 137.0989
Epoch 39/50
721/721 [=====] - 0s 389us/step - loss: 116.6981 -
val_loss: 136.3788
Epoch 40/50
721/721 [=====] - 0s 421us/step - loss: 115.8347 -
val_loss: 138.1166
Epoch 41/50
721/721 [=====] - 0s 688us/step - loss: 115.2098 -
val_loss: 138.9060
Epoch 42/50
721/721 [=====] - 0s 415us/step - loss: 114.5285 -
val_loss: 139.1216
Epoch 43/50
721/721 [=====] - 0s 416us/step - loss: 113.5577 -
val_loss: 140.4369
Epoch 44/50
721/721 [=====] - 0s 637us/step - loss: 112.9891 -
val_loss: 140.6825
Epoch 45/50
721/721 [=====] - 0s 411us/step - loss: 112.3483 -
val_loss: 141.5391
Epoch 46/50
721/721 [=====] - 0s 496us/step - loss: 111.5387 -
val_loss: 142.1855
Epoch 47/50
721/721 [=====] - 0s 419us/step - loss: 111.2179 -
val_loss: 140.9655
Epoch 48/50
721/721 [=====] - 0s 416us/step - loss: 110.5176 -
val_loss: 139.9818
Epoch 49/50
721/721 [=====] - 0s 410us/step - loss: 110.3294 -

```



```

val_loss: 143.7165
Epoch 50/50
721/721 [=====] - 0s 386us/step - loss: 109.2572 -
val_loss: 142.5937

```

[23]: <keras.callbacks.History at 0x7fee15c44390>

6 Evaluate the model on the test data

```

[24]: loss_val = model.evaluate(X_test, y_test)
      y_pred = model.predict(X_test)
      loss_val

```

```

309/309 [=====] - 0s 129us/step

```

[24]: 41407465.32038835

6.0.1 compute the mean squared error between the predicted concrete strength and the actual concrete strength.

Let's import the mean_squared_error function from Scikit-learn.

```

[25]: from sklearn.metrics import mean_squared_error

```

```

[26]: mean_square_error = mean_squared_error(y_test, y_pred)
      mean = np.mean(mean_square_error)
      standard_deviation = np.std(mean_square_error)
      print(mean, standard_deviation)

```

```

41407464.414271615 0.0

```

7 Create a list of 50 mean squared errors and report mean and the standard deviation of the mean squared errors.

```

[27]: total_mean_squared_errors = 50
      epochs = 50
      mean_squared_errors = []
      for i in range(0, total_mean_squared_errors):
          X_train, X_test, y_train, y_test = train_test_split(predictors, target,
          ↪test_size=0.3, random_state=i)
          model.fit(X_train, y_train, epochs=epochs, verbose=0)
          MSE = model.evaluate(X_test, y_test, verbose=0)
          print("MSE "+str(i+1)+": "+str(MSE))
          y_pred = model.predict(X_test)
          mean_square_error = mean_squared_error(y_test, y_pred)
          mean_squared_errors.append(mean_square_error)

```

```

mean_squared_errors = np.array(mean_squared_errors)
mean = np.mean(mean_squared_errors)
standard_deviation = np.std(mean_squared_errors)

print("Below is the mean and standard deviation of "
      ↳+str(total_mean_squared_errors) + " mean squared errors without normalized_
      ↳data. Total number of epochs for each training is: " +str(epochs) + "\n")
print("Mean: "+str(mean))
print("Standard Deviation: "+str(standard_deviation))

```

```

MSE 1: 2485.8806255056634
MSE 2: 364.48710365665767
MSE 3: 244.60458502414542
MSE 4: 270.6536970910131
MSE 5: 253.46854047559225
MSE 6: 215.75344305439674
MSE 7: 257.0338008349767
MSE 8: 184.53906190742566
MSE 9: 210.34155293189977
MSE 10: 183.46824670686692
MSE 11: 169.52499508163305
MSE 12: 164.02616916891054
MSE 13: 146.459690501389
MSE 14: 142.36362751093498
MSE 15: 100.41088844965962
MSE 16: 54.366279293418316
MSE 17: 55.93005685898864
MSE 18: 51.60272004457739
MSE 19: 41.15331311519092
MSE 20: 46.940897105195376
MSE 21: 43.17192497376871
MSE 22: 42.603018134157246
MSE 23: 39.24455702574893
MSE 24: 40.623416715455285
MSE 25: 45.03258888466844
MSE 26: 45.45151618537779
MSE 27: 45.71576132666332
MSE 28: 40.48105153451074
MSE 29: 46.40764478418048
MSE 30: 43.98078355511415
MSE 31: 45.48971606998382
MSE 32: 37.09484625325619
MSE 33: 41.74664067141832
MSE 34: 40.72372613061207
MSE 35: 37.425077888957894

```

MSE 36: 45.64947622725107
MSE 37: 46.295517183816166
MSE 38: 42.44215817744678
MSE 39: 38.2584443509
MSE 40: 39.126816610688145
MSE 41: 43.44314438162498
MSE 42: 37.77803346948716
MSE 43: 43.21549740886997
MSE 44: 41.35343080045336
MSE 45: 51.65417371830122
MSE 46: 43.734879070886905
MSE 47: 41.88216848280823
MSE 48: 43.21275382674628
MSE 49: 39.55350061064785
MSE 50: 44.282121417591874

Below is the mean and standard deviation of 50 mean squared errors without normalized data. Total number of epochs for each training is: 50

Mean: 138.40167421886906
Standard Deviation: 345.0033454170249

8 THANK YOU